

## *If statements/data selection/NA*

### *Introduction*

Please install the packages:

- ggplot2
- tufte

### *What is TRUE/FALSE?*

Today we will be working a lot with TRUE and FALSE.

Let us start off with a recap of last week. TRUE/FALSE are `boolean` variables. We can print them:

```
print(TRUE)

## [1] TRUE

print(FALSE)

## [1] FALSE
```

We can also assign them:

```
a <- TRUE
print(a)

## [1] TRUE

b <- FALSE
print(b)

## [1] FALSE
```

We can also **CREATE** them from questions:

```
3 == 4

## [1] FALSE
```

We can also save the result of the question (“is 3 equal to 4?”) in the variable `a` and then print the result:

```
a <- 3 == 4
a

## [1] FALSE
```

We can also ask other questions:

Note the use of `==` <br> Double `==`  
is a **QUESTION** <br> Single `=`  
is an **ASSIGNMENT**

**\*\*\*** means **NOT**

```

3 != 4 # is 3 not equal to 4?

## [1] TRUE

3 < 4 # is 3 less than 4?

## [1] TRUE

3 <= 4 # is 3 less than or equal to 4?

## [1] TRUE

3 > 4 # is 3 greater than 4?

## [1] FALSE

3 >= 4 # is 3 greater than or equal to 4?

## [1] FALSE

c(1, 2, 3, 4) == 3 # FOUR questions simultaneously

## [1] FALSE FALSE TRUE FALSE

c(3, 4) < c(4, 1) # TWO questions simultaneously (not recommended)

## [1] TRUE FALSE

```

The above questions (==, !=, <, <=, >, >=) **MUST** have either:

1. Have one single value on the right side of the question (recommended)
2. Have variables that are the same length on the left and the right of the question (not recommended)

The below questions can have different length variables on each side of the question:

```

3 %in% c(1, 2) # is 3 equal to 1 or 2?

## [1] FALSE

!3 %in% c(1, 2) # is 3 NOT equal to 1 or 2?

## [1] TRUE

c(1, 3) %in% c(2, 3, 4, 5) # TWO questions simultaneously

## [1] FALSE TRUE

```

We can obviously repeat all of these questions using variables instead of numbers:

The number of questions we ask **\*\*ALWAYS\*\*** corresponds to the length of the variable on the left side of the question

We do not use **\*\*c\*\*** as a variable because it is already a function

```

a <- 3
b <- 4
x <- c(1, 2)
a != b  # is 3 not equal to 4?

## [1] TRUE

a < b  # is 3 less than 4?

## [1] TRUE

a <= b  # is 3 less than or equal to 4?

## [1] TRUE

a > b  # is 3 greater than 4?

## [1] FALSE

a >= b  # is 3 greater than or equal to 4?

## [1] FALSE

a %in% x  # is 3 equal to 1 or 2?

## [1] FALSE

!a %in% x  # is 3 NOT equal to 1 or 2?

## [1] TRUE

```

### Exercises

Check to see if myAge is greater than or equal to ageLimitForBuyingAlcohol:

```

myAge <- 13
ageLimitForBuyingAlcohol <- 18

```

*# your code goes here*

Check to see if myFriendsAges are greater than or equal to ageLimitForBuyingAlcohol:

```

myFriendsAges <- c(15, 16, 16, 14, 20)
ageLimitForBuyingAlcohol <- 18

```

*# your code goes here*

Check to see if placeWhereILive exists in validCityNames:

```
placeWhereILive <- "Norway"
validCityNames <- c("Oslo", "Bergen", "Trondheim")
```

```
# your code goes here
```

Check to see if `placesWhereMyFriendsLive` are **NOT** in `placesInTheEU`:

```
placesWhereMyFriendsLive <- c("Australia", "UK",
                              "London", "Paris")
placesInTheEU <- c("France", "Paris", "Sweden")
```

```
# your code goes here
```

*AND/OR*

If we have multiple questions we can join them together using `&` (AND) and `|` (OR).

```
TRUE & TRUE
```

```
## [1] TRUE
```

```
TRUE & FALSE
```

```
## [1] FALSE
```

```
FALSE & FALSE
```

```
## [1] FALSE
```

```
TRUE | TRUE
```

```
## [1] TRUE
```

```
TRUE | FALSE
```

```
## [1] TRUE
```

```
FALSE | FALSE
```

```
## [1] FALSE
```

Put into more practical examples:

```
(3 < 4) & (2 < 5)
```

```
## [1] TRUE
```

```
(3 < 4) & (5 < 2)
```

Whenever using `|` and `&` it is always smart to use `()` to ensure your order of operations is correct

```
## [1] FALSE

(4 < 3) & (5 < 2)

## [1] FALSE

(3 < 4) | (2 < 5)

## [1] TRUE

(3 < 4) | (5 < 2)

## [1] TRUE

(4 < 3) | (5 < 2)

## [1] FALSE
```

## *Data Selection*

### *Vectors*

When we want to select data, we can either provide the location of the data we want (index) or a vector of TRUE/FALSE that essentially specifies `include/exclude` for every datapoint.

```
myData <- c("a", "b", "c", "d", "e")
myData

## [1] "a" "b" "c" "d" "e"

myData[2] # using indexes

## [1] "b"

myData[c(2, 4)] # using indexes

## [1] "b" "d"

myData[c(FALSE, TRUE, TRUE, FALSE, FALSE)] # using TRUE/FALSE

## [1] "b" "c"
```

### *data.frames*

We use a data.frame provided to us by the R package `ggplot2` (you might need to install this package):

```
data(diamonds, package = "ggplot2")
diamonds <- diamonds[1:20, ]
diamonds
```

```

##      carat      cut color clarity depth table
## 1  0.23     Ideal     E    SI2   61.5    55
## 2  0.21   Premium     E    SI1   59.8    61
## 3  0.23     Good     E    VS1   56.9    65
## 4  0.29   Premium     I    VS2   62.4    58
## 5  0.31     Good     J    SI2   63.3    58
## 6  0.24 Very Good     J   VVS2   62.8    57
## 7  0.24 Very Good     I   VVS1   62.3    57
## 8  0.26 Very Good     H    SI1   61.9    55
## 9  0.22     Fair     E    VS2   65.1    61
## 10 0.23 Very Good     H    VS1   59.4    61
## 11 0.30     Good     J    SI1   64.0    55
## 12 0.23     Ideal     J    VS1   62.8    56
## 13 0.22   Premium     F    SI1   60.4    61
## 14 0.31     Ideal     J    SI2   62.2    54
## 15 0.20   Premium     E    SI2   60.2    62
## 16 0.32   Premium     E     I1   60.9    58
## 17 0.30     Ideal     I    SI2   62.0    54
## 18 0.30     Good     J    SI1   63.4    54
## 19 0.30     Good     J    SI1   63.8    56
## 20 0.30 Very Good     J    SI1   62.7    59

##      price     x     y     z
## 1    326 3.95 3.98 2.43
## 2    326 3.89 3.84 2.31
## 3    327 4.05 4.07 2.31
## 4    334 4.20 4.23 2.63
## 5    335 4.34 4.35 2.75
## 6    336 3.94 3.96 2.48
## 7    336 3.95 3.98 2.47
## 8    337 4.07 4.11 2.53
## 9    337 3.87 3.78 2.49
## 10   338 4.00 4.05 2.39
## 11   339 4.25 4.28 2.73
## 12   340 3.93 3.90 2.46
## 13   342 3.88 3.84 2.33
## 14   344 4.35 4.37 2.71
## 15   345 3.79 3.75 2.27
## 16   345 4.38 4.42 2.68
## 17   348 4.31 4.34 2.68
## 18   351 4.23 4.29 2.70
## 19   351 4.23 4.26 2.71
## 20   351 4.21 4.27 2.66

```

Remember that each column of a data.frame is a vector, so we can

reuse what we learnt with vectors (selecting the first three values in the vector):

```
diamonds$carat[c(1:3)]

## [1] 0.23 0.21 0.23

diamonds$cut[c(1:3)]

## [1] Ideal   Premium Good
## 5 Levels: Fair < Good < ... < Ideal
```

We can also select the first three rows of a data.frame:

```
diamonds[c(1:3), ]

##   carat    cut color clarity depth table
## 1  0.23  Ideal     E    SI2   61.5     55
## 2  0.21 Premium     E    SI1   59.8     61
## 3  0.23   Good     E    VS1   56.9     65
##   price    x     y     z
## 1   326 3.95 3.98 2.43
## 2   326 3.89 3.84 2.31
## 3   327 4.05 4.07 2.31
```

We reference cells in data.frames by [ROWS, COLUMNS]. So diamonds[c(1:3),] means 'first three rows, all the columns'

The first two rows and the second and third columns (not recommended):

```
diamonds[c(1:2), c(2:3)]

##           cut color
## 1   Ideal     E
## 2 Premium     E
```

We can also specify the columns by name (recommended):

```
diamonds[c(1:2), c("cut", "color")]

##           cut color
## 1   Ideal     E
## 2 Premium     E
```

We can also use a vector of TRUE/FALSE instead of indexes:

```
diamonds[c(TRUE, TRUE, FALSE, FALSE, FALSE, FALSE,
  FALSE, FALSE, FALSE, FALSE, FALSE, FALSE,
  FALSE, FALSE, FALSE, FALSE, FALSE, FALSE,
  FALSE, FALSE), c("cut", "color")]
```

```
##      cut color
## 1   Ideal    E
## 2 Premium    E
```

**Most importantly** we can select rows by asking questions:

```
isCutPremium <- diamonds$cut == "Premium"
isCutPremium

## [1] FALSE  TRUE FALSE  TRUE FALSE FALSE
## [7] FALSE FALSE FALSE FALSE FALSE FALSE
## [13]  TRUE FALSE  TRUE  TRUE FALSE FALSE
## [19] FALSE FALSE

diamonds[isCutPremium, ]

##      carat      cut color clarity depth table
## 2    0.21 Premium    E    SI1  59.8    61
## 4    0.29 Premium    I    VS2  62.4    58
## 13   0.22 Premium    F    SI1  60.4    61
## 15   0.20 Premium    E    SI2  60.2    62
## 16   0.32 Premium    E    I1  60.9    58
##      price      x      y      z
## 2      326  3.89  3.84  2.31
## 4      334  4.20  4.23  2.63
## 13     342  3.88  3.84  2.33
## 15     345  3.79  3.75  2.27
## 16     345  4.38  4.42  2.68
```

And we can make this more elegant:

```
diamonds[diamonds$cut == "Premium", ]

##      carat      cut color clarity depth table
## 2    0.21 Premium    E    SI1  59.8    61
## 4    0.29 Premium    I    VS2  62.4    58
## 13   0.22 Premium    F    SI1  60.4    61
## 15   0.20 Premium    E    SI2  60.2    62
## 16   0.32 Premium    E    I1  60.9    58
##      price      x      y      z
## 2      326  3.89  3.84  2.31
## 4      334  4.20  4.23  2.63
## 13     342  3.88  3.84  2.33
## 15     345  3.79  3.75  2.27
## 16     345  4.38  4.42  2.68
```

We can also ask multiple questions:



```
diamonds[diamonds$cut == "Premium" & diamonds$color ==
  "E", ]
```

```
##      carat      cut color clarity depth table
## 2    0.21 Premium      E      SI1  59.8     61
## 15   0.20 Premium      E      SI2  60.2     62
## 16   0.32 Premium      E       I1  60.9     58
##      price      x      y      z
## 2      326  3.89  3.84  2.31
## 15     345  3.79  3.75  2.27
## 16     345  4.38  4.42  2.68
```

```
diamonds[diamonds$cut == "Premium" | diamonds$color ==
  "E", ]
```

```
##      carat      cut color clarity depth table
## 1    0.23   Ideal      E      SI2  61.5     55
## 2    0.21 Premium      E      SI1  59.8     61
## 3    0.23    Good      E      VS1  56.9     65
## 4    0.29 Premium      I      VS2  62.4     58
## 9    0.22   Fair      E      VS2  65.1     61
## 13   0.22 Premium      F      SI1  60.4     61
## 15   0.20 Premium      E      SI2  60.2     62
## 16   0.32 Premium      E       I1  60.9     58
##      price      x      y      z
## 1      326  3.95  3.98  2.43
## 2      326  3.89  3.84  2.31
## 3      327  4.05  4.07  2.31
## 4      334  4.20  4.23  2.63
## 9      337  3.87  3.78  2.49
## 13     342  3.88  3.84  2.33
## 15     345  3.79  3.75  2.27
## 16     345  4.38  4.42  2.68
```

```
diamonds[diamonds$cut %in% c("Premium", "Good") |
  diamonds$color == "E", ]
```

```
##      carat      cut color clarity depth table
## 1    0.23   Ideal      E      SI2  61.5     55
## 2    0.21 Premium      E      SI1  59.8     61
## 3    0.23    Good      E      VS1  56.9     65
## 4    0.29 Premium      I      VS2  62.4     58
## 5    0.31    Good      J      SI2  63.3     58
## 9    0.22   Fair      E      VS2  65.1     61
## 11   0.30    Good      J      SI1  64.0     55
## 13   0.22 Premium      F      SI1  60.4     61
```

```
## 15 0.20 Premium      E      SI2 60.2    62
## 16 0.32 Premium      E      I1 60.9    58
## 18 0.30      Good      J      SI1 63.4    54
## 19 0.30      Good      J      SI1 63.8    56
##      price      x      y      z
## 1      326 3.95 3.98 2.43
## 2      326 3.89 3.84 2.31
## 3      327 4.05 4.07 2.31
## 4      334 4.20 4.23 2.63
## 5      335 4.34 4.35 2.75
## 9      337 3.87 3.78 2.49
## 11     339 4.25 4.28 2.73
## 13     342 3.88 3.84 2.33
## 15     345 3.79 3.75 2.27
## 16     345 4.38 4.42 2.68
## 18     351 4.23 4.29 2.70
## 19     351 4.23 4.26 2.71
```

```
diamonds[diamonds$cut %in% c("Premium", "Good") &
  diamonds$price < 350, ]
```

```
##      carat      cut color clarity depth table
## 2    0.21 Premium      E      SI1 59.8    61
## 3    0.23      Good      E      VS1 56.9    65
## 4    0.29 Premium      I      VS2 62.4    58
## 5    0.31      Good      J      SI2 63.3    58
## 11   0.30      Good      J      SI1 64.0    55
## 13   0.22 Premium      F      SI1 60.4    61
## 15   0.20 Premium      E      SI2 60.2    62
## 16   0.32 Premium      E      I1 60.9    58
##      price      x      y      z
## 2      326 3.89 3.84 2.31
## 3      327 4.05 4.07 2.31
## 4      334 4.20 4.23 2.63
## 5      335 4.34 4.35 2.75
## 11     339 4.25 4.28 2.73
## 13     342 3.88 3.84 2.33
## 15     345 3.79 3.75 2.27
## 16     345 4.38 4.42 2.68
```

Once we have selected the rows we are interested in, we can then choose columns/variables:

```
diamonds[diamonds$cut == "Premium" & diamonds$color ==
  "E", ]$price
```

```
## [1] 326 345 345

diamonds[diamonds$cut == "Premium" | diamonds$color ==
  "E", ]$price

## [1] 326 326 327 334 337 342 345 345

diamonds[diamonds$cut %in% c("Premium", "Good") |
  diamonds$color == "E", ]$price

## [1] 326 326 327 334 335 337 339 342 345 345
## [11] 351 351

diamonds[diamonds$cut %in% c("Premium", "Good") &
  diamonds$price < 350, ]$price

## [1] 326 327 334 335 339 342 345 345
```

You can start to work with the data to get summary statistics:

```
mean(diamonds[diamonds$cut == "Premium" & diamonds$color ==
  "E", ]$price)

## [1] 338.6667

sd(diamonds[diamonds$cut == "Premium" & diamonds$color ==
  "E", ]$price)

## [1] 10.96966

quantile(diamonds[diamonds$cut == "Premium" &
  diamonds$color == "E", ]$price)

## 0% 25% 50% 75% 100%
## 326.0 335.5 345.0 345.0 345.0
```

Remember that behind each of these row selections is a vector containing TRUE/FALSE that includes/excludes certain rows:

```
diamonds$cut == "Premium" & diamonds$color ==
  "E"

## [1] FALSE TRUE FALSE FALSE FALSE FALSE
## [7] FALSE FALSE FALSE FALSE FALSE FALSE
## [13] FALSE FALSE TRUE TRUE FALSE FALSE
## [19] FALSE FALSE

diamonds$cut == "Premium" | diamonds$color ==
  "E"
```

```
## [1] TRUE TRUE TRUE TRUE FALSE FALSE
## [7] FALSE FALSE TRUE FALSE FALSE FALSE
## [13] TRUE FALSE TRUE TRUE FALSE FALSE
## [19] FALSE FALSE

diamonds$cut %in% c("Premium", "Good") | diamonds$color ==
  "E"

## [1] TRUE TRUE TRUE TRUE TRUE FALSE
## [7] FALSE FALSE TRUE FALSE TRUE FALSE
## [13] TRUE FALSE TRUE TRUE FALSE TRUE
## [19] TRUE FALSE

diamonds$cut %in% c("Premium", "Good") & diamonds$price <
  350

## [1] FALSE TRUE TRUE TRUE TRUE FALSE
## [7] FALSE FALSE FALSE FALSE TRUE FALSE
## [13] TRUE FALSE TRUE TRUE FALSE FALSE
## [19] FALSE FALSE
```

### *Data manipulation/cleaning*

Once you have selected your rows, you can also manipulate your data:

```
diamonds # before
```

	carat	cut	color	clarity	depth	table
## 1	0.23	Ideal	E	SI2	61.5	55
## 2	0.21	Premium	E	SI1	59.8	61
## 3	0.23	Good	E	VS1	56.9	65
## 4	0.29	Premium	I	VS2	62.4	58
## 5	0.31	Good	J	SI2	63.3	58
## 6	0.24	Very Good	J	VVS2	62.8	57
## 7	0.24	Very Good	I	VVS1	62.3	57
## 8	0.26	Very Good	H	SI1	61.9	55
## 9	0.22	Fair	E	VS2	65.1	61
## 10	0.23	Very Good	H	VS1	59.4	61
## 11	0.30	Good	J	SI1	64.0	55
## 12	0.23	Ideal	J	VS1	62.8	56
## 13	0.22	Premium	F	SI1	60.4	61
## 14	0.31	Ideal	J	SI2	62.2	54
## 15	0.20	Premium	E	SI2	60.2	62
## 16	0.32	Premium	E	I1	60.9	58
## 17	0.30	Ideal	I	SI2	62.0	54
## 18	0.30	Good	J	SI1	63.4	54
## 19	0.30	Good	J	SI1	63.8	56

```
## 20  0.30 Very Good      J      SI1  62.7    59
##      price      x      y      z
## 1    326 3.95 3.98 2.43
## 2    326 3.89 3.84 2.31
## 3    327 4.05 4.07 2.31
## 4    334 4.20 4.23 2.63
## 5    335 4.34 4.35 2.75
## 6    336 3.94 3.96 2.48
## 7    336 3.95 3.98 2.47
## 8    337 4.07 4.11 2.53
## 9    337 3.87 3.78 2.49
## 10   338 4.00 4.05 2.39
## 11   339 4.25 4.28 2.73
## 12   340 3.93 3.90 2.46
## 13   342 3.88 3.84 2.33
## 14   344 4.35 4.37 2.71
## 15   345 3.79 3.75 2.27
## 16   345 4.38 4.42 2.68
## 17   348 4.31 4.34 2.68
## 18   351 4.23 4.29 2.70
## 19   351 4.23 4.26 2.71
## 20   351 4.21 4.27 2.66
```

```
diamonds[diamonds$cut == "Premium" & diamonds$color ==
"E", ]$y <- 1000 # manipulate
diamonds # after
```

```
##      carat      cut color clarity depth table
## 1    0.23      Ideal      E      SI2   61.5    55
## 2    0.21    Premium      E      SI1   59.8    61
## 3    0.23       Good      E      VS1   56.9    65
## 4    0.29    Premium      I      VS2   62.4    58
## 5    0.31       Good      J      SI2   63.3    58
## 6    0.24 Very Good      J      VVS2   62.8    57
## 7    0.24 Very Good      I      VVS1   62.3    57
## 8    0.26 Very Good      H      SI1   61.9    55
## 9    0.22       Fair      E      VS2   65.1    61
## 10   0.23 Very Good      H      VS1   59.4    61
## 11   0.30       Good      J      SI1   64.0    55
## 12   0.23      Ideal      J      VS1   62.8    56
## 13   0.22    Premium      F      SI1   60.4    61
## 14   0.31      Ideal      J      SI2   62.2    54
## 15   0.20    Premium      E      SI2   60.2    62
## 16   0.32    Premium      E       I1   60.9    58
## 17   0.30      Ideal      I      SI2   62.0    54
```

```
## 18 0.30      Good      J      SI1 63.4   54
## 19 0.30      Good      J      SI1 63.8   56
## 20 0.30 Very Good      J      SI1 62.7   59
##      price      x          y      z
## 1      326 3.95      3.98 2.43
## 2      326 3.89 1000.00 2.31
## 3      327 4.05      4.07 2.31
## 4      334 4.20      4.23 2.63
## 5      335 4.34      4.35 2.75
## 6      336 3.94      3.96 2.48
## 7      336 3.95      3.98 2.47
## 8      337 4.07      4.11 2.53
## 9      337 3.87      3.78 2.49
## 10     338 4.00      4.05 2.39
## 11     339 4.25      4.28 2.73
## 12     340 3.93      3.90 2.46
## 13     342 3.88      3.84 2.33
## 14     344 4.35      4.37 2.71
## 15     345 3.79 1000.00 2.27
## 16     345 4.38 1000.00 2.68
## 17     348 4.31      4.34 2.68
## 18     351 4.23      4.29 2.70
## 19     351 4.23      4.26 2.71
## 20     351 4.21      4.27 2.66
```

This is how we clean our data. Here we change cut=="Premium" to Ideal:

```
diamonds # before
```

```
##      carat      cut color clarity depth table
## 1      0.23      Ideal      E      SI2 61.5   55
## 2      0.21    Premium      E      SI1 59.8   61
## 3      0.23      Good      E      VS1 56.9   65
## 4      0.29    Premium      I      VS2 62.4   58
## 5      0.31      Good      J      SI2 63.3   58
## 6      0.24 Very Good      J      VVS2 62.8   57
## 7      0.24 Very Good      I      VVS1 62.3   57
## 8      0.26 Very Good      H      SI1 61.9   55
## 9      0.22      Fair      E      VS2 65.1   61
## 10     0.23 Very Good      H      VS1 59.4   61
## 11     0.30      Good      J      SI1 64.0   55
## 12     0.23      Ideal      J      VS1 62.8   56
## 13     0.22    Premium      F      SI1 60.4   61
## 14     0.31      Ideal      J      SI2 62.2   54
```

```
## 15 0.20 Premium E SI2 60.2 62
## 16 0.32 Premium E I1 60.9 58
## 17 0.30 Ideal I SI2 62.0 54
## 18 0.30 Good J SI1 63.4 54
## 19 0.30 Good J SI1 63.8 56
## 20 0.30 Very Good J SI1 62.7 59
## price x y z
## 1 326 3.95 3.98 2.43
## 2 326 3.89 1000.00 2.31
## 3 327 4.05 4.07 2.31
## 4 334 4.20 4.23 2.63
## 5 335 4.34 4.35 2.75
## 6 336 3.94 3.96 2.48
## 7 336 3.95 3.98 2.47
## 8 337 4.07 4.11 2.53
## 9 337 3.87 3.78 2.49
## 10 338 4.00 4.05 2.39
## 11 339 4.25 4.28 2.73
## 12 340 3.93 3.90 2.46
## 13 342 3.88 3.84 2.33
## 14 344 4.35 4.37 2.71
## 15 345 3.79 1000.00 2.27
## 16 345 4.38 1000.00 2.68
## 17 348 4.31 4.34 2.68
## 18 351 4.23 4.29 2.70
## 19 351 4.23 4.26 2.71
## 20 351 4.21 4.27 2.66
```

```
diamonds[diamonds$cut == "Premium", ]$cut <- "Ideal" # manipulate
diamonds # after
```

```
## carat cut color clarity depth table
## 1 0.23 Ideal E SI2 61.5 55
## 2 0.21 Ideal E SI1 59.8 61
## 3 0.23 Good E VS1 56.9 65
## 4 0.29 Ideal I VS2 62.4 58
## 5 0.31 Good J SI2 63.3 58
## 6 0.24 Very Good J VVS2 62.8 57
## 7 0.24 Very Good I VVS1 62.3 57
## 8 0.26 Very Good H SI1 61.9 55
## 9 0.22 Fair E VS2 65.1 61
## 10 0.23 Very Good H VS1 59.4 61
## 11 0.30 Good J SI1 64.0 55
## 12 0.23 Ideal J VS1 62.8 56
## 13 0.22 Ideal F SI1 60.4 61
```

```
## 14 0.31      Ideal      J      SI2 62.2    54
## 15 0.20      Ideal      E      SI2 60.2    62
## 16 0.32      Ideal      E      I1  60.9    58
## 17 0.30      Ideal      I      SI2 62.0    54
## 18 0.30      Good       J      SI1 63.4    54
## 19 0.30      Good       J      SI1 63.8    56
## 20 0.30 Very Good      J      SI1 62.7    59
##      price      x      y      z
## 1      326 3.95      3.98 2.43
## 2      326 3.89 1000.00 2.31
## 3      327 4.05      4.07 2.31
## 4      334 4.20      4.23 2.63
## 5      335 4.34      4.35 2.75
## 6      336 3.94      3.96 2.48
## 7      336 3.95      3.98 2.47
## 8      337 4.07      4.11 2.53
## 9      337 3.87      3.78 2.49
## 10     338 4.00      4.05 2.39
## 11     339 4.25      4.28 2.73
## 12     340 3.93      3.90 2.46
## 13     342 3.88      3.84 2.33
## 14     344 4.35      4.37 2.71
## 15     345 3.79 1000.00 2.27
## 16     345 4.38 1000.00 2.68
## 17     348 4.31      4.34 2.68
## 18     351 4.23      4.29 2.70
## 19     351 4.23      4.26 2.71
## 20     351 4.21      4.27 2.66
```

We can also add different columns together:

```
diamonds # before
```

```
##      carat      cut color clarity depth table
## 1    0.23      Ideal      E      SI2 61.5    55
## 2    0.21      Ideal      E      SI1 59.8    61
## 3    0.23      Good       E      VS1 56.9    65
## 4    0.29      Ideal      I      VS2 62.4    58
## 5    0.31      Good       J      SI2 63.3    58
## 6    0.24 Very Good      J      VVS2 62.8    57
## 7    0.24 Very Good      I      VVS1 62.3    57
## 8    0.26 Very Good      H      SI1 61.9    55
## 9    0.22      Fair       E      VS2 65.1    61
## 10   0.23 Very Good      H      VS1 59.4    61
## 11   0.30      Good       J      SI1 64.0    55
```



```
## 12 0.23      Ideal      J      VS1 62.8 56
## 13 0.22      Ideal      F      SI1 60.4 61
## 14 0.31      Ideal      J      SI2 62.2 54
## 15 0.20      Ideal      E      SI2 60.2 62
## 16 0.32      Ideal      E      I1 60.9 58
## 17 0.30      Ideal      I      SI2 62.0 54
## 18 0.30      Good       J      SI1 63.4 54
## 19 0.30      Good       J      SI1 63.8 56
## 20 0.30 Very Good      J      SI1 62.7 59
```

```
##      price      x      y      z
## 1      326 3.95      3.98 2.43
## 2      326 3.89 1000.00 2.31
## 3      327 4.05      4.07 2.31
## 4      334 4.20      4.23 2.63
## 5      335 4.34      4.35 2.75
## 6      336 3.94      3.96 2.48
## 7      336 3.95      3.98 2.47
## 8      337 4.07      4.11 2.53
## 9      337 3.87      3.78 2.49
## 10     338 4.00      4.05 2.39
## 11     339 4.25      4.28 2.73
## 12     340 3.93      3.90 2.46
## 13     342 3.88      3.84 2.33
## 14     344 4.35      4.37 2.71
## 15     345 3.79 1000.00 2.27
## 16     345 4.38 1000.00 2.68
## 17     348 4.31      4.34 2.68
## 18     351 4.23      4.29 2.70
## 19     351 4.23      4.26 2.71
## 20     351 4.21      4.27 2.66
```

```
diamonds[diamonds$cut == "Ideal", ]$x <- diamonds[diamonds$cut ==
  "Ideal", ]$y + diamonds[diamonds$cut == "Ideal",
  ]$z
```

```
diamonds # after
```

```
##      carat      cut color clarity depth table
## 1      0.23      Ideal      E      SI2 61.5 55
## 2      0.21      Ideal      E      SI1 59.8 61
## 3      0.23      Good       E      VS1 56.9 65
## 4      0.29      Ideal      I      VS2 62.4 58
## 5      0.31      Good       J      SI2 63.3 58
## 6      0.24 Very Good      J      VVS2 62.8 57
## 7      0.24 Very Good      I      VVS1 62.3 57
## 8      0.26 Very Good      H      SI1 61.9 55
```

```
## 9 0.22 Fair E VS2 65.1 61
## 10 0.23 Very Good H VS1 59.4 61
## 11 0.30 Good J SI1 64.0 55
## 12 0.23 Ideal J VS1 62.8 56
## 13 0.22 Ideal F SI1 60.4 61
## 14 0.31 Ideal J SI2 62.2 54
## 15 0.20 Ideal E SI2 60.2 62
## 16 0.32 Ideal E I1 60.9 58
## 17 0.30 Ideal I SI2 62.0 54
## 18 0.30 Good J SI1 63.4 54
## 19 0.30 Good J SI1 63.8 56
## 20 0.30 Very Good J SI1 62.7 59
## price x y z
## 1 326 6.41 3.98 2.43
## 2 326 1002.31 1000.00 2.31
## 3 327 4.05 4.07 2.31
## 4 334 6.86 4.23 2.63
## 5 335 4.34 4.35 2.75
## 6 336 3.94 3.96 2.48
## 7 336 3.95 3.98 2.47
## 8 337 4.07 4.11 2.53
## 9 337 3.87 3.78 2.49
## 10 338 4.00 4.05 2.39
## 11 339 4.25 4.28 2.73
## 12 340 6.36 3.90 2.46
## 13 342 6.17 3.84 2.33
## 14 344 7.08 4.37 2.71
## 15 345 1002.27 1000.00 2.27
## 16 345 1002.68 1000.00 2.68
## 17 348 7.02 4.34 2.68
## 18 351 4.23 4.29 2.70
## 19 351 4.23 4.26 2.71
## 20 351 4.21 4.27 2.66
```

When we need to make multiple references to particular row selections, it is often cleaner to create a variable that contains the row selections:

```
# get row selection
rows <- diamonds$cut %in% c("Ideal", "Good") |
  diamonds$color == "E"
rows

## [1] TRUE TRUE TRUE TRUE TRUE FALSE
## [7] FALSE FALSE TRUE FALSE TRUE TRUE
```

```
## [13] TRUE TRUE TRUE TRUE TRUE TRUE
## [19] TRUE FALSE
```

```
diamonds # before
```

```
##      carat      cut color clarity depth table
## 1  0.23     Ideal     E    SI2   61.5    55
## 2  0.21     Ideal     E    SI1   59.8    61
## 3  0.23     Good      E    VS1   56.9    65
## 4  0.29     Ideal     I    VS2   62.4    58
## 5  0.31     Good      J    SI2   63.3    58
## 6  0.24 Very Good     J    VVS2   62.8    57
## 7  0.24 Very Good     I    VVS1   62.3    57
## 8  0.26 Very Good     H    SI1   61.9    55
## 9  0.22     Fair      E    VS2   65.1    61
## 10 0.23 Very Good     H    VS1   59.4    61
## 11 0.30     Good      J    SI1   64.0    55
## 12 0.23     Ideal     J    VS1   62.8    56
## 13 0.22     Ideal     F    SI1   60.4    61
## 14 0.31     Ideal     J    SI2   62.2    54
## 15 0.20     Ideal     E    SI2   60.2    62
## 16 0.32     Ideal     E     I1   60.9    58
## 17 0.30     Ideal     I    SI2   62.0    54
## 18 0.30     Good      J    SI1   63.4    54
## 19 0.30     Good      J    SI1   63.8    56
## 20 0.30 Very Good     J    SI1   62.7    59

##      price      x      y      z
## 1    326    6.41    3.98 2.43
## 2    326 1002.31 1000.00 2.31
## 3    327    4.05    4.07 2.31
## 4    334    6.86    4.23 2.63
## 5    335    4.34    4.35 2.75
## 6    336    3.94    3.96 2.48
## 7    336    3.95    3.98 2.47
## 8    337    4.07    4.11 2.53
## 9    337    3.87    3.78 2.49
## 10   338    4.00    4.05 2.39
## 11   339    4.25    4.28 2.73
## 12   340    6.36    3.90 2.46
## 13   342    6.17    3.84 2.33
## 14   344    7.08    4.37 2.71
## 15   345 1002.27 1000.00 2.27
## 16   345 1002.68 1000.00 2.68
## 17   348    7.02    4.34 2.68
## 18   351    4.23    4.29 2.70
```

```
## 19 351 4.23 4.26 2.71
## 20 351 4.21 4.27 2.66
```

```
diamonds[rows, ]$x <- diamonds[rows, ]$y + diamonds[rows,
  ]$z * 100
```

```
diamonds # after
```

```
##      carat      cut color clarity depth table
## 1  0.23    Ideal     E    SI2   61.5    55
## 2  0.21    Ideal     E    SI1   59.8    61
## 3  0.23     Good     E    VS1   56.9    65
## 4  0.29    Ideal     I    VS2   62.4    58
## 5  0.31     Good     J    SI2   63.3    58
## 6  0.24 Very Good     J   VVS2   62.8    57
## 7  0.24 Very Good     I   VVS1   62.3    57
## 8  0.26 Very Good     H    SI1   61.9    55
## 9  0.22     Fair     E    VS2   65.1    61
## 10 0.23 Very Good     H    VS1   59.4    61
## 11 0.30     Good     J    SI1   64.0    55
## 12 0.23    Ideal     J    VS1   62.8    56
## 13 0.22    Ideal     F    SI1   60.4    61
## 14 0.31    Ideal     J    SI2   62.2    54
## 15 0.20    Ideal     E    SI2   60.2    62
## 16 0.32    Ideal     E     I1   60.9    58
## 17 0.30    Ideal     I    SI2   62.0    54
## 18 0.30     Good     J    SI1   63.4    54
## 19 0.30     Good     J    SI1   63.8    56
## 20 0.30 Very Good     J    SI1   62.7    59

##      price      x      y      z
## 1    326 246.98   3.98 2.43
## 2    326 1231.00 1000.00 2.31
## 3    327 235.07   4.07 2.31
## 4    334 267.23   4.23 2.63
## 5    335 279.35   4.35 2.75
## 6    336   3.94   3.96 2.48
## 7    336   3.95   3.98 2.47
## 8    337   4.07   4.11 2.53
## 9    337 252.78   3.78 2.49
## 10   338   4.00   4.05 2.39
## 11   339 277.28   4.28 2.73
## 12   340 249.90   3.90 2.46
## 13   342 236.84   3.84 2.33
## 14   344 275.37   4.37 2.71
## 15   345 1227.00 1000.00 2.27
## 16   345 1268.00 1000.00 2.68
```

```
## 17    348  272.34    4.34 2.68
## 18    351  274.29    4.29 2.70
## 19    351  275.26    4.26 2.71
## 20    351    4.21    4.27 2.66
```

### *Creating new data.frames*

Once you have selected your rows, you can also save it to a new data.frame:

```
myNewDataFrame <- diamonds[diamonds$cut == "Ideal" &
  diamonds$color == "E", ]
myNewDataFrame
```

```
##      carat   cut color clarity depth table
## 1   0.23 Ideal     E    SI2  61.5    55
## 2   0.21 Ideal     E    SI1  59.8    61
## 15  0.20 Ideal     E    SI2  60.2    62
## 16  0.32 Ideal     E     I1  60.9    58
##      price      x      y      z
## 1     326  246.98   3.98 2.43
## 2     326 1231.00 1000.00 2.31
## 15    345 1227.00 1000.00 2.27
## 16    345 1268.00 1000.00 2.68
```

### *Exercises*

**Task 1:** Select all rows with:

- colour equals E or I AND
- price less than 400

Solution 1:

```
diamonds[diamonds$color %in% c("E", "I") & diamonds$price <
  400, ]
```

```
##      carat      cut color clarity depth table
## 1   0.23    Ideal     E    SI2  61.5    55
## 2   0.21    Ideal     E    SI1  59.8    61
## 3   0.23     Good     E    VS1  56.9    65
## 4   0.29    Ideal     I    VS2  62.4    58
## 7   0.24 Very Good     I   VVS1  62.3    57
## 9   0.22     Fair     E    VS2  65.1    61
## 15  0.20    Ideal     E    SI2  60.2    62
## 16  0.32    Ideal     E     I1  60.9    58
## 17  0.30    Ideal     I    SI2  62.0    54
```

```
##      price      x      y      z
## 1      326 246.98    3.98 2.43
## 2      326 1231.00 1000.00 2.31
## 3      327 235.07    4.07 2.31
## 4      334 267.23    4.23 2.63
## 7      336    3.95    3.98 2.47
## 9      337 252.78    3.78 2.49
## 15     345 1227.00 1000.00 2.27
## 16     345 1268.00 1000.00 2.68
## 17     348 272.34    4.34 2.68
```

**Task 2:** Select all rows with:

- depth less than 63
- price more than 300

Solution 2:

```
# your code goes here
```

**Task 3:**

Set z to 400 for all rows with:

- cut not Ideal
- price more than 300

Solution 3:

```
# your code goes here
```

### *Control flow (if statements)*

Today's module has so far only focused on using TRUE/FALSE to select data.

We can also use TRUE/FALSE to change the flow of the program.

```
if (TRUE) {
  print("a")
} else {
  print("b")
}

## [1] "a"
```

Above we have used TRUE/FALSE to select which lines of code would run. This is called an if statement.

These will be used more frequently in further modules, but at this stage you could use them in conjunction with flags at the start of your script.

```

analyseCheapDiamonds <- TRUE # a flag

if (analyseCheapDiamonds) {
  analysisData <- diamonds[diamonds$price <
    350, ]
} else {
  analysisData <- diamonds
}

analysisData

##      carat      cut color clarity depth table
## 1  0.23    Ideal     E    SI2   61.5    55
## 2  0.21    Ideal     E    SI1   59.8    61
## 3  0.23     Good     E    VS1   56.9    65
## 4  0.29    Ideal     I    VS2   62.4    58
## 5  0.31     Good     J    SI2   63.3    58
## 6  0.24 Very Good     J   VVS2   62.8    57
## 7  0.24 Very Good     I   VVS1   62.3    57
## 8  0.26 Very Good     H    SI1   61.9    55
## 9  0.22     Fair     E    VS2   65.1    61
## 10 0.23 Very Good     H    VS1   59.4    61
## 11 0.30     Good     J    SI1   64.0    55
## 12 0.23    Ideal     J    VS1   62.8    56
## 13 0.22    Ideal     F    SI1   60.4    61
## 14 0.31    Ideal     J    SI2   62.2    54
## 15 0.20    Ideal     E    SI2   60.2    62
## 16 0.32    Ideal     E     I1   60.9    58
## 17 0.30    Ideal     I    SI2   62.0    54
##      price      x      y      z
## 1    326  246.98   3.98  2.43
## 2    326 1231.00 1000.00  2.31
## 3    327   235.07   4.07  2.31
## 4    334   267.23   4.23  2.63
## 5    335   279.35   4.35  2.75
## 6    336    3.94   3.96  2.48
## 7    336    3.95   3.98  2.47
## 8    337    4.07   4.11  2.53
## 9    337   252.78   3.78  2.49
## 10   338    4.00   4.05  2.39
## 11   339   277.28   4.28  2.73
## 12   340   249.90   3.90  2.46
## 13   342   236.84   3.84  2.33
## 14   344   275.37   4.37  2.71

```

```
## 15    345 1227.00 1000.00 2.27
## 16    345 1268.00 1000.00 2.68
## 17    348  272.34    4.34 2.68
```

In the above script the user can change the top line `analyseCheapDiamonds` `<- TRUE` to be `TRUE/FALSE` if they want `analysisData` to be only cheap diamonds or all of the diamonds. The rest of the script would continue using `analysisData`, which would mean that you could write 1 analysis script that could be used for 2 (or more) different datasets.

We can also write more complicated if statements. In the following code, the flag `analysis` can take the values: `main`, `sensitivity1`, `sensitivity2` and changes the dataset accordingly:

```
analysis <- "main"

if (analysis == "main") {
  analysisData <- diamonds
} else if (analysis == "sensitivity1") {
  analysisData <- diamonds[diamonds$cut != "Good",
]
} else if (analysis == "sensitivity2") {
  analysisData <- diamonds[diamonds$color ==
    "E", ]
} else {
  stop("not a valid analysis!")
}
```

`analysisData`

```
##      carat      cut color clarity depth table
## 1    0.23    Ideal     E     SI2   61.5    55
## 2    0.21    Ideal     E     SI1   59.8    61
## 3    0.23     Good     E     VS1   56.9    65
## 4    0.29    Ideal     I     VS2   62.4    58
## 5    0.31     Good     J     SI2   63.3    58
## 6    0.24 Very Good     J    VVS2   62.8    57
## 7    0.24 Very Good     I    VVS1   62.3    57
## 8    0.26 Very Good     H     SI1   61.9    55
## 9    0.22     Fair     E     VS2   65.1    61
## 10   0.23 Very Good     H     VS1   59.4    61
## 11   0.30     Good     J     SI1   64.0    55
## 12   0.23    Ideal     J     VS1   62.8    56
## 13   0.22    Ideal     F     SI1   60.4    61
## 14   0.31    Ideal     J     SI2   62.2    54
## 15   0.20    Ideal     E     SI2   60.2    62
```



```
## 16 0.32      Ideal      E      I1 60.9    58
## 17 0.30      Ideal      I      SI2 62.0    54
## 18 0.30      Good      J      SI1 63.4    54
## 19 0.30      Good      J      SI1 63.8    56
## 20 0.30 Very Good      J      SI1 62.7    59
##      price      x      y      z
## 1      326 246.98    3.98 2.43
## 2      326 1231.00 1000.00 2.31
## 3      327 235.07    4.07 2.31
## 4      334 267.23    4.23 2.63
## 5      335 279.35    4.35 2.75
## 6      336      3.94    3.96 2.48
## 7      336      3.95    3.98 2.47
## 8      337      4.07    4.11 2.53
## 9      337 252.78    3.78 2.49
## 10     338      4.00    4.05 2.39
## 11     339 277.28    4.28 2.73
## 12     340 249.90    3.90 2.46
## 13     342 236.84    3.84 2.33
## 14     344 275.37    4.37 2.71
## 15     345 1227.00 1000.00 2.27
## 16     345 1268.00 1000.00 2.68
## 17     348 272.34    4.34 2.68
## 18     351 274.29    4.29 2.70
## 19     351 275.26    4.26 2.71
## 20     351      4.21    4.27 2.66
```

If you were running regression analyses, then you could take it even further and have different outcomes, exposures, and confounders for each analysis:

```
analysis <- "main"

if (analysis == "main") {
  analysisData <- diamonds
  outcome <- "price"
  exposure <- "carat"
  confounders <- c("color", "clarity", "depth")
} else if (analysis == "sensitivity1") {
  analysisData <- diamonds[diamonds$cut != "Good",
    ]
  outcome <- "price"
  exposure <- "carat"
  confounders <- c("cut", "clarity")
} else if (analysis == "sensitivity2") {
```

```

analysisData <- diamonds[diamonds$color ==
  "E", ]
outcome <- "price"
exposure <- "cut"
confounders <- c("depth")
} else {
  stop("not a valid analysis!")
}

```

```
# create new data.frame
```

```

analysisData <- analysisData[, c(outcome, exposure,
  confounders)]
analysisData

```

```

##      price carat color clarity depth
## 1      326  0.23     E     SI2   61.5
## 2      326  0.21     E     SI1   59.8
## 3      327  0.23     E     VS1   56.9
## 4      334  0.29     I     VS2   62.4
## 5      335  0.31     J     SI2   63.3
## 6      336  0.24     J    VVS2   62.8
## 7      336  0.24     I    VVS1   62.3
## 8      337  0.26     H     SI1   61.9
## 9      337  0.22     E     VS2   65.1
## 10     338  0.23     H     VS1   59.4
## 11     339  0.30     J     SI1   64.0
## 12     340  0.23     J     VS1   62.8
## 13     342  0.22     F     SI1   60.4
## 14     344  0.31     J     SI2   62.2
## 15     345  0.20     E     SI2   60.2
## 16     345  0.32     E      I1   60.9
## 17     348  0.30     I     SI2   62.0
## 18     351  0.30     J     SI1   63.4
## 19     351  0.30     J     SI1   63.8
## 20     351  0.30     J     SI1   62.7

```

```
# rename variables in the data.frame
```

```

names(analysisData)[c(1:2)] <- c("outcome", "exposure")
names(analysisData)[c(3:ncol(analysisData))] <- paste0("confounder",
  1:length(confounders))
analysisData

```

```

##      outcome exposure confounder1 confounder2
## 1      326      0.23           E          SI2
## 2      326      0.21           E          SI1

```

```
## 3      327      0.23      E      VS1
## 4      334      0.29      I      VS2
## 5      335      0.31      J      SI2
## 6      336      0.24      J      VVS2
## 7      336      0.24      I      VVS1
## 8      337      0.26      H      SI1
## 9      337      0.22      E      VS2
## 10     338      0.23      H      VS1
## 11     339      0.30      J      SI1
## 12     340      0.23      J      VS1
## 13     342      0.22      F      SI1
## 14     344      0.31      J      SI2
## 15     345      0.20      E      SI2
## 16     345      0.32      E      I1
## 17     348      0.30      I      SI2
## 18     351      0.30      J      SI1
## 19     351      0.30      J      SI1
## 20     351      0.30      J      SI1
```

```
##      confounder3
```

```
## 1      61.5
## 2      59.8
## 3      56.9
## 4      62.4
## 5      63.3
## 6      62.8
## 7      62.3
## 8      61.9
## 9      65.1
## 10     59.4
## 11     64.0
## 12     62.8
## 13     60.4
## 14     62.2
## 15     60.2
## 16     60.9
## 17     62.0
## 18     63.4
## 19     63.8
## 20     62.7
```

```
# 'analyse' some of the data.frame
```

```
mean(analysisData$outcome)
```

```
## [1] 339.4
```


```
mean(analysisData$exposure)
```

```
## [1] 0.262

cor(analysisData$outcome, analysisData$exposure)

## [1] 0.5250715

plot(outcome ~ exposure, data = analysisData)
```



### *NA/Missing/Invalid Data*

There are two kinds of missing/invalid data:

- What **YOU** know is missing/invalid
- What **R** knows is missing/invalid

```
myWeightOverTime <- c(80, 86, 76, 80, -9, -9,
  -9, 0, 0, 0, 89)
myWeightOverTime

## [1] 80 86 76 80 -9 -9 -9 0 0 0 89

mean(myWeightOverTime)

## [1] 34.90909
```

Here, it is obvious to **me** that the values between 70 and 90 are valid, while -9 and 0 are obviously missing data.

For **R** all of these values are valid and real. The **ONLY** value that **R** considers to be missing/invalid is **NA**. So we need to translate my knowledge of the data, and my understanding of what values are missing/invalid into **NAs** that **R** can understand:

```
myWeightOverTime[(myWeightOverTime < 70) | (myWeightOverTime >
  90)] <- NA
myWeightOverTime

## [1] 80 86 76 80 NA NA NA NA NA NA 89
```

We can now analyse the data:

```
mean(myWeightOverTime)

## [1] NA
```

This doesn't work, because there are NAs in the data. We need to explicitly tell R to ignore the NAs:

```
mean(myWeightOverTime, na.rm = TRUE)

## [1] 82.2
```

Of course, all of this applies to data.frames:

```
diamonds # before

##   carat    cut color clarity depth table
## 1  0.23   Ideal    E   SI2   61.5    55
## 2  0.21   Ideal    E   SI1   59.8    61
## 3  0.23    Good    E   VS1   56.9    65
## 4  0.29   Ideal    I   VS2   62.4    58
## 5  0.31    Good    J   SI2   63.3    58
## 6  0.24 Very Good    J  VVS2   62.8    57
## 7  0.24 Very Good    I  VVS1   62.3    57
## 8  0.26 Very Good    H   SI1   61.9    55
## 9  0.22    Fair    E   VS2   65.1    61
## 10 0.23 Very Good    H   VS1   59.4    61
## 11 0.30    Good    J   SI1   64.0    55
## 12 0.23   Ideal    J   VS1   62.8    56
## 13 0.22   Ideal    F   SI1   60.4    61
## 14 0.31   Ideal    J   SI2   62.2    54
## 15 0.20   Ideal    E   SI2   60.2    62
## 16 0.32   Ideal    E    I1   60.9    58
## 17 0.30   Ideal    I   SI2   62.0    54
## 18 0.30    Good    J   SI1   63.4    54
## 19 0.30    Good    J   SI1   63.8    56
## 20 0.30 Very Good    J   SI1   62.7    59
##   price      x      y      z
## 1   326  246.98   3.98  2.43
## 2   326 1231.00 1000.00  2.31
```

```
## 3    327  235.07    4.07 2.31
## 4    334  267.23    4.23 2.63
## 5    335  279.35    4.35 2.75
## 6    336    3.94    3.96 2.48
## 7    336    3.95    3.98 2.47
## 8    337    4.07    4.11 2.53
## 9    337  252.78    3.78 2.49
## 10   338    4.00    4.05 2.39
## 11   339  277.28    4.28 2.73
## 12   340  249.90    3.90 2.46
## 13   342  236.84    3.84 2.33
## 14   344  275.37    4.37 2.71
## 15   345 1227.00 1000.00 2.27
## 16   345 1268.00 1000.00 2.68
## 17   348  272.34    4.34 2.68
## 18   351  274.29    4.29 2.70
## 19   351  275.26    4.26 2.71
## 20   351    4.21    4.27 2.66
```

```
diamonds[diamonds$cut %in% c("Good", "Fair"),
]$price <- NA # manipulation
diamonds # after
```

```
##      carat      cut color clarity depth table
## 1    0.23    Ideal     E     SI2   61.5    55
## 2    0.21    Ideal     E     SI1   59.8    61
## 3    0.23     Good     E     VS1   56.9    65
## 4    0.29    Ideal     I     VS2   62.4    58
## 5    0.31     Good     J     SI2   63.3    58
## 6    0.24 Very Good     J    VVS2   62.8    57
## 7    0.24 Very Good     I    VVS1   62.3    57
## 8    0.26 Very Good     H     SI1   61.9    55
## 9    0.22     Fair     E     VS2   65.1    61
## 10   0.23 Very Good     H     VS1   59.4    61
## 11   0.30     Good     J     SI1   64.0    55
## 12   0.23    Ideal     J     VS1   62.8    56
## 13   0.22    Ideal     F     SI1   60.4    61
## 14   0.31    Ideal     J     SI2   62.2    54
## 15   0.20    Ideal     E     SI2   60.2    62
## 16   0.32    Ideal     E      I1   60.9    58
## 17   0.30    Ideal     I     SI2   62.0    54
## 18   0.30     Good     J     SI1   63.4    54
## 19   0.30     Good     J     SI1   63.8    56
## 20   0.30 Very Good     J     SI1   62.7    59
##      price      x      y      z
```

```
## 1    326  246.98    3.98 2.43
## 2    326 1231.00 1000.00 2.31
## 3     NA  235.07    4.07 2.31
## 4    334  267.23    4.23 2.63
## 5     NA  279.35    4.35 2.75
## 6    336    3.94    3.96 2.48
## 7    336    3.95    3.98 2.47
## 8    337    4.07    4.11 2.53
## 9     NA  252.78    3.78 2.49
## 10   338    4.00    4.05 2.39
## 11   NA  277.28    4.28 2.73
## 12   340  249.90    3.90 2.46
## 13   342  236.84    3.84 2.33
## 14   344  275.37    4.37 2.71
## 15   345 1227.00 1000.00 2.27
## 16   345 1268.00 1000.00 2.68
## 17   348  272.34    4.34 2.68
## 18   NA  274.29    4.29 2.70
## 19   NA  275.26    4.26 2.71
## 20   351    4.21    4.27 2.66
```

```
mean(diamonds$price)
```

```
## [1] NA
```

```
mean(diamonds$price, na.rm = TRUE)
```

```
## [1] 339.1429
```

We can now introduce the function (question) `is.na`:

```
is.na(4)
```

```
## [1] FALSE
```

```
is.na("hello")
```

```
## [1] FALSE
```

```
is.na(NA)
```

```
## [1] TRUE
```

And if we apply it to a vector that contains NAs:

```
diamonds$price
```

```
## [1] 326 326 NA 334 NA 336 336 337 NA 338
## [11] NA 340 342 344 345 345 348 NA NA 351
```

```
is.na(diamonds$price)

## [1] FALSE FALSE TRUE FALSE TRUE FALSE
## [7] FALSE FALSE TRUE FALSE TRUE FALSE
## [13] FALSE FALSE FALSE FALSE FALSE TRUE
## [19] TRUE FALSE
```

So we can then use `is.na` in row selections, the same as all of the other questions (`=`, `!=`, `<`, `<=`, `>`, `>=`, `%in%`):

```
diamonds[is.na(diamonds$price), ]

##   carat  cut color clarity depth table
## 3  0.23 Good   E     VS1   56.9     65
## 5  0.31 Good   J     SI2   63.3     58
## 9  0.22 Fair   E     VS2   65.1     61
## 11 0.30 Good   J     SI1   64.0     55
## 18 0.30 Good   J     SI1   63.4     54
## 19 0.30 Good   J     SI1   63.8     56
##   price      x      y      z
## 3      NA 235.07 4.07 2.31
## 5      NA 279.35 4.35 2.75
## 9      NA 252.78 3.78 2.49
## 11     NA 277.28 4.28 2.73
## 18     NA 274.29 4.29 2.70
## 19     NA 275.26 4.26 2.71
```

```
diamonds[!is.na(diamonds$price), ]

##   carat      cut color clarity depth table
## 1  0.23    Ideal   E     SI2   61.5     55
## 2  0.21    Ideal   E     SI1   59.8     61
## 4  0.29    Ideal   I     VS2   62.4     58
## 6  0.24 Very Good   J     VVS2   62.8     57
## 7  0.24 Very Good   I     VVS1   62.3     57
## 8  0.26 Very Good   H     SI1   61.9     55
## 10 0.23 Very Good   H     VS1   59.4     61
## 12 0.23    Ideal   J     VS1   62.8     56
## 13 0.22    Ideal   F     SI1   60.4     61
## 14 0.31    Ideal   J     SI2   62.2     54
## 15 0.20    Ideal   E     SI2   60.2     62
## 16 0.32    Ideal   E     I1   60.9     58
## 17 0.30    Ideal   I     SI2   62.0     54
## 20 0.30 Very Good   J     SI1   62.7     59
##   price      x      y      z
## 1    326 246.98  3.98 2.43
```



```
## 2      326 1231.00 1000.00 2.31
## 4      334  267.23    4.23 2.63
## 6      336    3.94    3.96 2.48
## 7      336    3.95    3.98 2.47
## 8      337    4.07    4.11 2.53
## 10     338    4.00    4.05 2.39
## 12     340  249.90    3.90 2.46
## 13     342  236.84    3.84 2.33
## 14     344  275.37    4.37 2.71
## 15     345 1227.00 1000.00 2.27
## 16     345 1268.00 1000.00 2.68
## 17     348  272.34    4.34 2.68
## 20     351    4.21    4.27 2.66
```

### *Sum and Mean on TRUE/FALSE*

You can use `sum` to count how many observations are true:

```
sum(c(TRUE, TRUE, FALSE, FALSE, FALSE, NA), na.rm = T)
```

```
## [1] 2
```

```
sum(diamonds$price < 350, na.rm = T)
```

```
## [1] 13
```

You can use `mean` to see the proportion of values that are true:

```
mean(c(TRUE, TRUE, FALSE, FALSE, FALSE, NA), na.rm = T)
```

```
## [1] 0.4
```

```
mean(diamonds$price < 350, na.rm = T)
```

```
## [1] 0.9285714
```